




Tackling inter-service RDF communication bottlenecks in the Nanopublication network with Jelly

Piotr Sowiński¹, Tobias Kuhn², Karolina Bogacka¹

¹ NeverBlink

² Knowledge Pixels



Google invented Protobuf to solve it...


But **5%** of their datacenter CPU cycles are still spent on ser/des!

Can your serialization keep up with the rest of the system?



Nanopublication network

Anatomy of a nanopublication



- FAIR by design
- 1 nanopub = 1 RDF dataset
(4 named graphs)
- ~50–200 triples
- Lots of them!

Nanopublication

Participation in: 2025-eu.semantics.cc P1344 ▾


I (Piotr Sowiński) participated in 2025-eu.semantics.cc .

The assertion above is attributed to Piotr Sowiński .


This nanopublication is created by me (Piotr Sowiński) .

```
13 ✓ sub:Head {  
14   ✓ this: np:hasAssertion sub:assertion;  
15     np:hasProvenance sub:provenance;  
16     np:hasPublicationInfo sub:pubinfo;  
17     a np:Nanopublication .  
18 }  
19  
20 ✓ sub:assertion {  
21   orcid:0000-0002-2543-9461 <http://www.wikidata.org/entity/P1344> <https://2025-eu.semantics.cc/> .  
22 }  
23  
24 ✓ sub:provenance {  
25   sub:assertion prov:wasAttributedTo orcid:0000-0002-2543-9461 .  
26 }  
27  
28 ✓ sub:pubinfo {  
29   orcid:0000-0002-2543-9461 foaf:name "Piotr Sowiński" .  
30  
31 ✓ this: dct:created "2025-08-31T09:12:56.973Z"^^xsd:dateTime;  
32   dct:creator orcid:0000-0002-2543-9461;  
33   dct:license <https://creativecommons.org/licenses/by/4.0/>;  
34   npx:wasCreatedAt <https://nanodash.knowledgepixels.com/>;  
35   rdfs:label "Participation in: 2025-eu.semantics.cc";  
36   nt:wasCreatedFromProvenanceTemplate <https://w3id.org/np/RA71Sq6MuK_TIC6JMSHvLtee3lpLoZD0qLJCLXevnrPoU>;  
37 ✓ nt:wasCreatedFromPubinfoTemplate <https://w3id.org/np/RA0J4vUn_dekg-U1kK3A0Et02p9mT2W003uGxLDec1jLw>,  
38   | <https://w3id.org/np/RAukAcWHRDlkqxk7H2XNSegc1WnHI569INvNr-xdptDGI>;  
39   nt:wasCreatedFromTemplate <https://w3id.org/np/RA580k5zFLCd9N7nPPrJgwURUTTgP2mkb2vg-4LBd0etpE> .  
40  
41 ✓ sub:sig npx:hasAlgorithm "RSA";  
42   npx:hasPublicKey "MIGfMA0GCSqGSIb3DQEBAQUAA4GNADCBiQKBgQCNMXM2Ib2J9WEFG510mfII9CoT6BURjAtQK8vpbdXJLC+WXTu3p/7U08mq24zKpiZNVa  
43   npx:hasSignature "ZS/S/ObM2dN0wtoXTFfkp5IUv1KYaktUZ85QD0QieqtCV07TJGZRzRO/UWjw6qad0tH91vt3fedf/2AnGxy09K8pPN0tU22/95L1/VD9qf  
44   npx:hasSignatureTarget this:;  
45   npx:signedBy orcid:0000-0002-2543-9461 .  
46 }  
47
```


Nanopublication network



Nanopublication network



Nanopublication network



Starting situation

- HTML / JSON list pages with links to individual nanopubs
- Individual nanopubs served as TriG files
- Accessing 60k nanopubs = 60k+ HTTP requests

Latest Nanopubs List (max. 1000)

1. [RAeNcV9gE7](https://w3id.org/np/RAeNcV9gE7rRHF5KuVeFu60B67IyHwcVqVuJ43wKe1X5)
2. [RABdZf6gri](https://w3id.org/np/RABdZf6gri)
3. [RAHeq2_sF4](https://w3id.org/np/RAHeq2_sF4)
4. [RAD2m1CfxA](https://w3id.org/np/RAd2m1CfxA)
5. [RA1JyVWpyV](https://w3id.org/np/RA1JyVWpyV)
6. [RA0ohu7SuQ](https://w3id.org/np/RA0ohu7SuQ)
7. [RAiEW_g7ws](https://w3id.org/np/RAiEW_g7ws)
8. [RAmcsRu2MU](https://w3id.org/np/RAmcsRu2MU)
9. [RAfdMk3PtG](https://w3id.org/np/RAfdMk3PtG)
10. [RA4-Qyqu2X](https://w3id.org/np/RA4-Qyqu2X)
11. [RAcId5yDwr](https://w3id.org/np/RAcId5yDwr)
12. [RAM4gTJg3C](https://w3id.org/np/RAM4gTJg3C)
13. [RAcIYMb12p](https://w3id.org/np/RAcIYMb12p)
14. [RA4eX94wB-](https://w3id.org/np/RA4eX94wB-)
15. [RAEoHdKtgv](https://w3id.org/np/RAEoHdKtgv)
16. [RAXFsDo32E](https://w3id.org/np/RAXFsDo32E)
17. [RA2mrYXX9](https://w3id.org/np/RA2mrYXX9)
18. [RAnIo0cT1k](https://w3id.org/np/RAnIo0cT1k)
19. [RAP-73pUtM](https://w3id.org/np/RAP-73pUtM)
20. [RAVaM_WPGG](https://w3id.org/np/RAVaM_WPGG)
21. [RAMdi28Cp8](https://w3id.org/np/RAMdi28Cp8)
22. [RAz1p4D6m-](https://w3id.org/np/RAz1p4D6m-)
23. [RAh60f2hQS](https://w3id.org/np/RAh60f2hQS)
24. [RA04i14Zsb](https://w3id.org/np/RA04i14Zsb)
25. [RAP41TrW7T](https://w3id.org/np/RAp41TrW7T)
26. [RA3iVIom0S](https://w3id.org/np/RA3iVIom0S)
27. [RA0kRa006n](https://w3id.org/np/RA0kRa006n)

Nanopublication

< Home

ID

<https://w3id.org/np/RAeNcV9gE7rRHF5KuVeFu60B67IyHwcVqVuJ43wKe1X5>

Formats

.trig | .trig.txt | .jelly | .jelly.txt | .jsonld | .jsonld.txt | .nq | .nq.txt | .xml | .xml.txt

Content

```
@prefix this: <https://w3id.org/np/RAeNcV9gE7rRHF5KuVeFu60B67IyHwcVqVuJ43wKe1X5> .
@prefix sub: <https://w3id.org/np/RAeNcV9gE7rRHF5KuVeFu60B67IyHwcVqVuJ43wKe1X5#sub> .
@prefix np: <http://www.nanopub.org/nschema#> .
@prefix dct: <http://purl.org/dc/terms/> .
@prefix nt: <https://w3id.org/np/o/ntemplate/> .
@prefix npx: <http://purl.org/nanopub/x/> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix orcid: <https://orcid.org/> .
@prefix prov: <http://www.w3.org/ns/prov#> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .

sub:Head {
    this: a np:Nanopublication;
    np:hasAssertion sub:assertion;
    np:hasProvenance sub:provenance;
```

Starting situation – issues

- TriG format is very slow to parse
- Repeated HTTP requests add a lot of overhead

The result:


- Very slow replication throughput
- Additional latency (1 round-trip for list, then 1 for nanopub)



Solution: Jelly

Jelly in a nutshell


- Binary RDF format based on Protobuf
- 100% open spec & open source (<https://w3id.org/jelly>)
- Very fast to write (**2x** faster than N-Triples in Jena)
- Very, very fast to read (**12x** faster than N-Triples)
- Reasonably well-compressed (**6x** smaller than N-Triples)



Jelly in a nutshell


Works with:

- Java (Apache Jena, RDF4J, Titanium)
- Python (RDFLib or no library)
- Rust (*Sophia*) – experimental, community-led
- Neo4j
- CLI application



How does Jelly work?

- Lightweight streaming compression algorithm
- For n triples:
 - $O(1)$ memory complexity
 - $O(n)$ time complexity
- Max supported triple count = ∞
- 1 file can contain 1 RDF document (graph or dataset)...
- ...or 1 file can contain **many** RDF documents (!)



Clients

**Nanopub Registry
API app (Java)**

**Nanopub Registry
DB (MongoDB)**

Nanopub 1

Jelly blob

Nanopub 2

Jelly blob

Nanopub 3

Jelly blob

Nanopub 4

Jelly blob


Nanopub 5

Jelly blob

Clients “Give me nanopub number 4, please!”

**Nanopub Registry
API app (Java)**

**Nanopub Registry
DB (MongoDB)**







Clients

*“Give me nanopub
number 4, please!”*

Nanopub Registry
API app (Java)

Nanopub Registry
DB (MongoDB)



Clients

“Give me nanopubs by Tobias Kuhn, please!”

Nanopub Registry
API app (Java)

Nanopub Registry
DB (MongoDB)

Nanopub 1

Jelly blob

Nanopub 2

Jelly blob

Nanopub 3

Jelly blob

Nanopub 4

Jelly blob

Nanopub 5

Jelly blob

Clients

“Give me nanopubs by Tobias Kuhn, please!”

HTTP response
(streaming)

Nanopub Registry
API app (Java)

Merged Jelly stream (1, 3, 4)

Transcoding

Jelly blob 1

Jelly blob 3

Jelly blob 4

Nanopub Registry
DB (MongoDB)

Nanopub 1

Jelly blob

Nanopub 2

Jelly blob

Nanopub 3

Jelly blob

Nanopub 4

Jelly blob

Nanopub 5

Jelly blob

Results

Naïve comparison: original

Takes >3 hours
to complete.



```
1 import json
2 import requests
3 import rdflib
4
5 list_response = requests.get(
6     'https://registry.knowledgepixels.com/nanopubs.json'
7 ).content
8 list_json = json.loads(list_response)
9 for i, item in enumerate(list_json):
10     if i % 1000 == 0:
11         print(f'Processed {i} nanopubs')
12     url = f'https://registry.knowledgepixels.com/np/{item}'
13     try:
14         response = requests.get(url, headers={
15             'Accept': 'application/trig'
16         })
17         g = rdflib.Dataset()
18         g.parse(source=response.content, format='trig')
19     except Exception as e:
20         print(f'Error retrieving nanopub from {url}: {e}')
21
```

Naïve comparison: Jelly

Takes ~4 seconds
to complete.

```
piotr@perun:~$ time wget https://registry.petapico.org/nanopubs.jelly -q -O- |  
jelly-cli rdf inspect  
stream_options:  
  stream_name: ""  
  physical_type: QUADS (2)  
  generalized_statements: false  
  rdf_star: false  
  max_name_table_size: 4000  
  max_prefix_table_size: 150  
  max_datatype_table_size: 32  
  logical_type: DATASETS (4)  
version: 2  
  
frames:  
  frame_count: 64925  
  row_count: 2467391  
  option_count: 1  
  triple_count: 0  
  quad_count: 1604486  
  graph_start_count: 0  
  graph_end_count: 0  
  namespace_count: 647906  
  name_count: 121096  
  prefix_count: 93892  
  datatype_count: 10  
  
real    0m4.196s  
user    0m0.314s  
sys     0m0.069s  
piotr@perun:~$ █
```

Raw ser/des throughput comparison (no HTTP overhead)



Platform: Oracle GraalVM 24+36.1, RDF4J 5.1.4, Jelly-JVM 2.10.3, Ryzen 9 7900 5.0 GHz, 64 GB RAM
Dataset: 10M nanopublications (RiverBench: nanopubs)

Why not pipelining, parallelization, caching...?

- More complex = more costly
- Hidden resource usage – overhead still largely exists!
- Depends on the client to “do things right”
- Jelly can also compress across nanopublication boundaries
- Caches don’t help!
 - Cache is usually completely cold

Conclusion

knowledge
pixels

- Communication went **from a bottleneck to a non-issue** thanks to Jelly
- **Live** on the nanopublication network: <https://nanopub.net>
- **Large potential for transferability:**
 - Mature tooling & documentation
 - **Use cases:** client-server communication, microservices, database dumps, streaming ingest, database replication, and more...
 - **Open community – anyone can contribute and use Jelly!**
- 100% open-source

<https://w3id.org/jelly>



★ Star us on
GitHub!

Backup slides

Solution summary

- Registry serves arbitrary subsets of nanopubs as a single streaming HTTP response
- Query & Registry consume the stream, unpack it, and process each nanopub individually

To retrieve 60k nanopubs:

- Original: 60k+ requests
- Jelly: **exactly 1 request**

Size comparison

RiverBench dataset: nanopubs, obtained with Apache Jena 5.1.0

